

Starting the Hive Service:

```
[hduser@localhost ~]$ hive --service metastore
ls: cannot access /usr/local/spark/lib/spark-assembly-*.jar: No such file or directory
Starting Hive Metastore Server
22/07/04 15:16:03 WARN conf.HiveConf: HiveConf of name hive.mapjoin.optimized.keys does not exist
22/07/04 15:16:03 WARN conf.HiveConf: HiveConf of name hive.optimize.multigroupby.common.distincts does not exist
22/07/04 15:16:03 WARN conf.HiveConf: HiveConf of name hive.mapjoin.lazy.hashtable does not exist
22/07/04 15:16:03 WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.min.worker.threads does not exist
22/07/04 15:16:03 WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.max.worker.threads does not exist
22/07/04 15:16:03 WARN conf.HiveConf: HiveConf of name hive.server2.logging.operation.verbose does not exist
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/phoenix-4.11.0-HBase-0.98-client.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Loading class `com.mysql.jdbc.Driver'. This is deprecated. The new driver class is `com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the
SPI and manual loading of the driver class is generally unnecessary.
```

Connecting to beeline:

```
[hduser@localhost ~]$ beeline
ls: cannot access /usr/local/spark/lib/spark-assembly-*.jar: No such file or directory
Beeline version 1.2.2 by Apache Hive
beeline> !connect jdbc:hive2://
SLF4J: Class path contains multiple SLF4J bindings.
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/phoenix-4.11.0-HBase-0.98-client.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hbase/lib/slf4j-log4j12-1.6.4.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: Found binding in [jar:file:/usr/local/hadoop/share/hadoop/common/lib/slf4j-log4j12-1.7.10.jar!/org/slf4j/impl/StaticLoggerBinder.class]
SLF4J: See http://www.slf4j.org/codes.html#multiple_bindings for an explanation.
SLF4J: Actual binding is of type [org.slf4j.impl.Log4jLoggerFactory]
Connecting to jdbc:hive2://
Enter username for jdbc:hive2://: scott
Enter password for jdbc:hive2://: *****
22/07/04 15:18:12 [main]: WARN conf.HiveConf: HiveConf of name hive.mapjoin.optimized.keys does not exist
22/07/04 15:18:12 [main]: WARN conf.HiveConf: HiveConf of name hive.optimize.multigroupby.common.distincts does not exist
22/07/04 15:18:12 [main]: WARN conf.HiveConf: HiveConf of name hive.mapjoin.lazy.hashtable does not exist
22/07/04 15:18:12 [main]: WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.min.worker.threads does not exist
22/07/04 15:18:12 [main]: WARN conf.HiveConf: HiveConf of name hive.server2.thrift.http.max.worker.threads does not exist
22/07/04 15:18:12 [main]: WARN conf.HiveConf: HiveConf of name hive.server2.logging.operation.verbose does not exist
22/07/04 15:18:12 [main]: WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Connected to Apache Hive (version 1.2.2)
Driver: Hive JDBC (version 1.2.2)
Transaction isolation: TRANSACTION_REPEATABLE_READ
0: jdbc:hive2://>
```

Static Partitioning

Creating Files:

```
[hduser@localhost ~]$ gedit Hyd.txt
[hduser@localhost ~]$ gedit Ban.txt
[hduser@localhost ~]$ cat Hyd.txt
1,Rahul,20,14000
2,Ramesh,30,14500
3,Guru,35,17850
[hduser@localhost ~]$ gedit Ban.txt
[hduser@localhost ~]$ cat Ban.txt
4,Giri,23,13890
5,Mohan,45,65000
5,Venu,40,45000
[hduser@localhost ~]$ gedit Delhi.txt
[hduser@localhost ~]$ cat Delhi.txt
7,Aravind,45,80000
8,Mukesh,50,45600
```

Moving to HDFS:

```
[hduser@localhost ~]$ hdfs dfs -put Hyd.txt /user/hduser/hive_practice
22/07/04 15:56:35 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[hduser@localhost ~]$ hdfs dfs -put Ban.txt /user/hduser/hive_practice
22/07/04 15:56:45 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[hduser@localhost ~]$ hdfs dfs -put Delhi.txt /user/hduser/hive_practice
22/07/04 15:56:55 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[hduser@localhost ~]$ hdfs dfs -ls /user/hduser/hive_practice/
22/07/04 15:57:02 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r-- 1 hduser supergroup 49 2022-07-04 15:56 /user/hduser/hive_practice/Ban.txt
-rw-r--r-- 1 hduser supergroup 37 2022-07-04 15:56 /user/hduser/hive_practice/Delhi.txt
-rw-r--r-- 1 hduser supergroup 51 2022-07-04 15:56 /user/hduser/hive_practice/Hyd.txt
[hduser@localhost ~]$
```

Table Creation:

```
0: jdbc:hive2://> create table Static_Partition_Demo(eid int,ename string,eage int) partitioned by(eloc string) row format delimited fields terminated by ','
lines terminated by '\n';
OK
No rows affected (7.046 seconds)
```

```
0: jdbc:hive2://> desc Static_Partition_Demo;
OK
22/07/04 16:01:30 [main]: WARN lazy.LazyStruct: Extra bytes detected at the end of the row! Ignoring similar problems.
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| eid      | int      |         |
| ename    | string   |         |
| eage     | int      |         |
| eloc     | string   |         |
| # Partition Information | NULL      | NULL    |
| # col_name | data_type | comment |
|          | NULL      | NULL    |
| eloc     | string   |         |
+-----+-----+-----+
9 rows selected (2.249 seconds)
```

Loading the Data using partitions:

```
0: jdbc:hive2://> load data inpath "/user/hduser/hive_practice/Hyd.txt" into table Static_Partition_Demo partition(eloc = "Hyd");
Loading data to table practice.static_partition_demo partition (eloc=Hyd)
Partition practice.static_partition_demo(eloc=Hyd) stats: [numFiles=1, numRows=0, totalSize=51, rawDataSize=0]
OK
No rows affected (3.609 seconds)
0: jdbc:hive2://> load data inpath "/user/hduser/hive_practice/Ban.txt" into table Static_Partition_Demo partition(eloc = "Ban");
Loading data to table practice.static_partition_demo partition (eloc=Ban)
Partition practice.static_partition_demo(eloc=Ban) stats: [numFiles=1, numRows=0, totalSize=49, rawDataSize=0]
OK
No rows affected (1.322 seconds)
0: jdbc:hive2://> load data inpath "/user/hduser/hive_practice/Delhi.txt" into table Static_Partition_Demo partition(eloc = "Delhi");
Loading data to table practice.static_partition_demo partition (eloc=Delhi)
Partition practice.static_partition_demo(eloc=Delhi) stats: [numFiles=1, numRows=0, totalSize=37, rawDataSize=0]
OK
No rows affected (1.473 seconds)
0: jdbc:hive2://>
```

```
22/07/04 16:05:51 [main]: WARN lazy.LazyStruct: Extra bytes detected at the end of the row! Ignoring similar problems.
+-----+-----+-----+-----+
| static_partition_demo.eid | static_partition_demo.ename | static_partition_demo.eage | static_partition_demo.eloc |
+-----+-----+-----+-----+
| 4 | Giri | 23 | Ban |
| 5 | Mohan | 45 | Ban |
| 5 | Venu | 40 | Ban |
| 7 | Aravind | 45 | Delhi |
| 8 | Mukesh | 50 | Delhi |
| 1 | Rahul | 20 | Hyd |
| 2 | Ramesh | 30 | Hyd |
| 3 | Guru | 35 | Hyd |
+-----+-----+-----+-----+
8 rows selected (1.749 seconds)
```

```
0: jdbc:hive2://> show partitions Static_Partition_Demo;
OK
+-----+
| partition |
+-----+
| eloc=Ban |
| eloc=Delhi |
| eloc=Hyd |
+-----+
3 rows selected (0.4 seconds)
```

```
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/static_partition_demo
22/07/04 16:09:02 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:05 /user/hive/warehouse/practice.db/static_partition_demo/eloc=Ban
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:05 /user/hive/warehouse/practice.db/static_partition_demo/eloc=Delhi
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:04 /user/hive/warehouse/practice.db/static_partition_demo/eloc=Hyd
```

```
0: jdbc:hive2://> select * from Static_Partition_Demo where eloc = "Hyd";
OK
22/07/04 16:10:31 [main]: WARN lazy.LazyStruct: Extra bytes detected at the end of the row! Ignoring similar problems.
+-----+-----+-----+-----+
| static_partition_demo.eid | static_partition_demo.ename | static_partition_demo.eage | static_partition_demo.eloc |
+-----+-----+-----+-----+
| 1 | Rahul | 20 | Hyd |
| 2 | Ramesh | 30 | Hyd |
| 3 | Guru | 35 | Hyd |
+-----+-----+-----+-----+
3 rows selected (1.221 seconds)
```

Dynamic Partitioning

```
[hduser@localhost ~]$ gedit Customers.txt
[hduser@localhost ~]$ cat Customers.txt
1,Kumar,23,Hyderabad
2,Manesh,34,Mumbai
3,Suresh,24,Hyderabad
4,Gopi,45,Pune
5,Surender,45,Pune
6,Krishna,56,Hyderabad
[hduser@localhost ~]$ hdfs dfs -ls /user/hduser/hive_practice
22/07/04 16:17:29 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[hduser@localhost ~]$ hdfs dfs -put Customers.txt /user/hduser/hive_practice
22/07/04 16:17:35 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
[hduser@localhost ~]$ hdfs dfs -ls /user/hduser/hive_practice
22/07/04 16:17:42 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 hduser supergroup 119 2022-07-04 16:17 /user/hduser/hive_practice/Customers.txt
[hduser@localhost ~]$
```

Creating table without Partitions and loading the data:

```
0: jdbc:hive2://> create table customer_nopart(cid int,cname string,cage int,cloc string) row format delimited fields terminated by ',' lines terminated by '\n';
OK
No rows affected (0.318 seconds)
0: jdbc:hive2://> load data inpath '/user/hduser/hive_practice/Customers.txt' into table customer_nopart;
Loading data to table practice.customer_nopart
Table practice.customer_nopart stats: [numFiles=1, totalSize=119]
OK
No rows affected (0.764 seconds)
0: jdbc:hive2://> select * from customer_nopart;
OK
+-----+-----+-----+-----+
| customer_nopart.cid | customer_nopart.cname | customer_nopart.cage | customer_nopart.cloc |
+-----+-----+-----+-----+
| 1 | Kumar | 23 | Hyderabad |
| 2 | Mahesh | 34 | Mumbai |
| 3 | Suresh | 24 | Hyderabad |
| 4 | Gopi | 45 | Pune |
| 5 | Surender | 45 | Pune |
| 6 | Krishna | 56 | Hyderabad |
+-----+-----+-----+-----+
6 rows selected (0.418 seconds)
```

Creating a table with Partition Column:

```
0: jdbc:hive2://> create table customer_part(cid int,cname string,cage int) partitioned by(cloc string) row format delimited fields terminated by ',' lines terminated by '\n';
OK
No rows affected (0.586 seconds)
0: jdbc:hive2://> desc customer_part;
OK
22/07/04 16:19:17 [main]: WARN LazyLazyStruct: Extra bytes detected at the end of the row! Ignoring similar problems.
+-----+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+-----+
| cid | int | |
| cname | string | |
| cage | int | |
| cloc | string | |
| # Partition Information | NULL | NULL |
| # col_name | data_type | comment |
| cloc | string | NULL |
+-----+-----+-----+-----+
0 rows selected (0.316 seconds)
0: jdbc:hive2://>
```

Inserting data from previous table to partition containing table:

```
hive> set hive.exec.dynamic.partition = true;
hive> set hive.exec.dynamic.partition.mode = nonstrict;
hive> insert into customer_part partition(cloc) select * from customer_nopart;
Query ID = hduser_20220704162917_40513f4b-7557-4092-a446-c9f844008a63
Total jobs = 3
Launching Job 1 out of 3
Number of reduce tasks is set to 0 since there's no reduce operator
Loading class 'com.mysql.jdbc.Driver'. This is deprecated. The new driver class is 'com.mysql.cj.jdbc.Driver'. The driver is automatically registered via the SPI and manual loading of the driver class is generally unnecessary.
Starting Job = job_1656927935583_0003, Tracking URL = http://inceptor:8088/proxy/application_1656927935583_0003/
Kill Command = /usr/local/hadoop/bin/hadoop job -kill job_1656927935583_0003
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 0
```

```
hive> select * from customer_part;
OK
1 Kumar 23 Hyderabad
3 Suresh 24 Hyderabad
6 Krishna 56 Hyderabad
2 Mahesh 34 Mumbai
4 Gopi 45 Pune
5 Surender 45 Pune
Time taken: 0.578 seconds, Fetched: 6 row(s)
hive>
```

```
hive> show partitions customer_part;
OK
cloc=Hyderabad
cloc=Mumbai
cloc=Pune
Time taken: 0.403 seconds, Fetched: 3 row(s)
hive>
```

```
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/customer_part
22/07/04 16:32:15 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:30 /user/hive/warehouse/practice.db/customer_part/cloc=Hyderabad
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:30 /user/hive/warehouse/practice.db/customer_part/cloc=Mumbai
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:30 /user/hive/warehouse/practice.db/customer_part/cloc=Pune
[hduser@localhost ~]$
```

```
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/customer_part/cloc=Pune
22/07/04 16:32:39 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 hduser supergroup 24 2022-07-04 16:30 /user/hive/warehouse/practice.db/customer_part/cloc=Pune/000000_0
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/customer_part/cloc=Hyderabad
22/07/04 16:32:48 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 hduser supergroup 36 2022-07-04 16:30 /user/hive/warehouse/practice.db/customer_part/cloc=Hyderabad/000000_0
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/customer_part/cloc=Mumbai
22/07/04 16:32:56 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 1 items
-rw-r--r-- 1 hduser supergroup 12 2022-07-04 16:30 /user/hive/warehouse/practice.db/customer_part/cloc=Mumbai/000000_0
[hduser@localhost ~]$
```

Bucketing

Creating table with some data without buckets:

```
0: jdbc:hive2://> create table customer_nopart(cid int,cname string,cage int,cloc string) row format delimited fields terminated by ',' lines terminated by '\n';
OK
No rows affected (0.318 seconds)
0: jdbc:hive2://> load data inpath '/user/hduser/hive_practice/Customers.txt' into table customer_nopart;
Loading data to table practice.customer_nopart
Table practice.customer_nopart stats: [numFiles=1, totalSize=119]
OK
No rows affected (0.764 seconds)
0: jdbc:hive2://> select * from customer_nopart;
OK
+-----+-----+-----+-----+
| customer_nopart.cid | customer_nopart.cname | customer_nopart.cage | customer_nopart.cloc |
+-----+-----+-----+-----+
| 1 | Kumar | 23 | Hyderabad |
| 2 | Mahesh | 34 | Mumbai |
| 3 | Suresh | 24 | Hyderabad |
| 4 | Gopi | 45 | Pune |
| 5 | Surrender | 45 | Pune |
| 6 | Krishna | 56 | Hyderabad |
+-----+-----+-----+-----+
6 rows selected (0.418 seconds)
```

Table creation with buckets:

```
hive> create table customer_buck(cid int,cname string,cage int,cloc string) clustered by(cloc) into 3 buckets row format delimited fields terminated by ',' lines terminated by '\n';
OK
Time taken: 0.635 seconds
hive> desc customer_buck;
OK
cid                int
cname              string
cage               int
cloc               string
Time taken: 0.25 seconds, Fetched: 4 row(s)
hive>
```

Loading the data into new table with buckets:

```
hive> insert into customer_buck select * from customer_nopart;
Query ID = hduser_20220704163606_d8ff9093-9bb0-4a70-8442-bf8120d08727
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
Starting Job = job_1656927935583_0004, Tracking URL = http://Inceptez:8088/proxy/application_1656927935583_0004/
Kill command = /usr/local/hadoop/bin/hadoop job -kill job_1656927935583_0004
Hadoop job information for Stage-1: number of mappers: 1; number of reducers: 3
2022-07-04 16:36:46,787 Stage-1 map = 0%, reduce = 0%
2022-07-04 16:37:00,405 Stage-1 map = 100%, reduce = 0%, Cumulative CPU 1.75 sec
2022-07-04 16:37:28,827 Stage-1 map = 100%, reduce = 44%, Cumulative CPU 3.33 sec
2022-07-04 16:37:32,597 Stage-1 map = 100%, reduce = 56%, Cumulative CPU 4.84 sec
2022-07-04 16:37:35,124 Stage-1 map = 100%, reduce = 67%, Cumulative CPU 5.94 sec
2022-07-04 16:37:37,295 Stage-1 map = 100%, reduce = 100%, Cumulative CPU 7.5 sec
```

```
hive> select * from customer_buck;
OK
5      Surrender      45      Pune
4      Gopi      45      Pune
2      Mahesh      34      Mumbai
6      Krishna      56      Hyderabad
3      Suresh      24      Hyderabad
1      Kumar      23      Hyderabad
Time taken: 0.483 seconds, Fetched: 6 row(s)
hive>
```

```
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/customer_buck
22/07/04 16:38:46 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r-- 1 hduser supergroup 34 2022-07-04 16:37 /user/hive/warehouse/practice.db/customer_buck/000000_0
-rw-r--r-- 1 hduser supergroup 19 2022-07-04 16:37 /user/hive/warehouse/practice.db/customer_buck/000001_0
-rw-r--r-- 1 hduser supergroup 66 2022-07-04 16:37 /user/hive/warehouse/practice.db/customer_buck/000002_0
[hduser@localhost ~]$
```

Partition with Bucketing:

Creating table with some data without buckets and partitions:

```
0: jdbc:hive2://> create table customer_nopart(cid int,cname string,cage int,cloc string) row format delimited fields terminated by ',' lines terminated by '\n';
OK
No rows affected (0.318 seconds)
0: jdbc:hive2://> load data inpath '/user/hduser/hive_practice/Customers.txt' into table customer_nopart;
Loading data to table practice.customer_nopart
Table practice.customer_nopart stats: [numFiles=1, totalSize=119]
OK
No rows affected (0.764 seconds)
0: jdbc:hive2://> select * from customer_nopart;
OK
+-----+-----+-----+-----+
| customer_nopart.cid | customer_nopart.cname | customer_nopart.cage | customer_nopart.cloc |
+-----+-----+-----+-----+
| 1 | Kumar | 23 | Hyderabad |
| 2 | Mahesh | 34 | Mumbai |
| 3 | Suresh | 24 | Hyderabad |
| 4 | Gopi | 45 | Pune |
| 5 | Surrender | 45 | Pune |
| 6 | Krishna | 56 | Hyderabad |
+-----+-----+-----+-----+
6 rows selected (0.418 seconds)
```

Creation of table with bucket and partition:

```
hive> create table customer_partbuck(cid int,cname string,cage int) partitioned by(cloc string) clustered by(cid) into 3 buckets row format delimited fields terminated by ',' lines terminated by '\n';
OK
Time taken: 0.423 seconds
hive> desc customer_partbuck;
OK
cid                int
cname              string
cage               int
cloc               string

# Partition Information
# col_name          data_type          comment
cloc               string

Time taken: 0.245 seconds, Fetched: 9 row(s)
hive>
```

Loading the data:

```
hive> insert into customer_partbuck partition(cloc) select * from customer_nopart;
Query ID = hduser_20220704164419_7e874635-fc66-480d-94b5-2e1b1bee41ac
Total jobs = 1
Launching Job 1 out of 1
Number of reduce tasks determined at compile time: 3
In order to change the average load for a reducer (in bytes):
  set hive.exec.reducers.bytes.per.reducer=<number>
In order to limit the maximum number of reducers:
  set hive.exec.reducers.max=<number>
In order to set a constant number of reducers:
  set mapreduce.job.reduces=<number>
```

```
hive> show partitions customer_partbuck;
OK
cloc=Hyderabad
cloc=Mumbai
cloc=Pune
Time taken: 0.471 seconds, Fetched: 3 row(s)
hive>
```

```
hive> select * from customer_partbuck;
OK
6      Krishna 56      Hyderabad
3      Suresh  24      Hyderabad
1      Kumar   23      Hyderabad
2      Mahesh  34      Mumbai
4      Gopi    45      Pune
5      Surender 45      Pune
Time taken: 0.638 seconds, Fetched: 6 row(s)
hive>
```

```
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/
22/07/04 16:47:08 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 9 items
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:37 /user/hive/warehouse/practice.db/customer_buck
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:18 /user/hive/warehouse/practice.db/customer_nopart
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:30 /user/hive/warehouse/practice.db/customer_part
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:46 /user/hive/warehouse/practice.db/customer_partbuck
drwxr-xr-x - hduser supergroup 0 2022-06-30 11:40 /user/hive/warehouse/practice.db/employee
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:05 /user/hive/warehouse/practice.db/static_partition_demo
drwxr-xr-x - hduser supergroup 0 2022-07-01 11:34 /user/hive/warehouse/practice.db/structdemo
drwxr-xr-x - hduser supergroup 0 2022-07-01 11:37 /user/hive/warehouse/practice.db/structdemo2
drwxr-xr-x - hduser supergroup 0 2022-07-01 10:27 /user/hive/warehouse/practice.db/studentsandmarks
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/customer_partbuck
22/07/04 16:47:20 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:46 /user/hive/warehouse/practice.db/customer_partbuck/cloc=Hyderabad
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:46 /user/hive/warehouse/practice.db/customer_partbuck/cloc=Mumbai
drwxr-xr-x - hduser supergroup 0 2022-07-04 16:46 /user/hive/warehouse/practice.db/customer_partbuck/cloc=Pune
[hduser@localhost ~]$ hdfs dfs -ls /user/hive/warehouse/practice.db/customer_partbuck/cloc=Pune
22/07/04 16:47:28 WARN util.NativeCodeLoader: Unable to load native-hadoop library for your platform... using builtin-java classes where applicable
Found 3 items
-rw-r--r-- 1 hduser supergroup 0 2022-07-04 16:46 /user/hive/warehouse/practice.db/customer_partbuck/cloc=Pune/000000_0
-rw-r--r-- 1 hduser supergroup 10 2022-07-04 16:46 /user/hive/warehouse/practice.db/customer_partbuck/cloc=Pune/000001_0
-rw-r--r-- 1 hduser supergroup 14 2022-07-04 16:46 /user/hive/warehouse/practice.db/customer_partbuck/cloc=Pune/000002_0
[hduser@localhost ~]$
```